

Test results:

E2E functional test:

1. Created a cloud function that can easily change returned status code.
(wewu_less/handlers/tester.py)
2. Added job using **WeWu API** with POST:

JavaScript

```
{
  "serviceUrl": "https://us-east1-wewu-410223.cloudfunctions.net/wewu_tester",
  "geoRegions": ["us-east1-b"],
  "primaryAdmin": {"email": "*****@gmail.com"},
  "secondaryAdmin": {"email": "*****@gmail.com"},
  "pollFrequencySecs": 5,
  "alertingWindowNumberOfCalls": 10,
  "alertingWindowCallsFailCount": 6,
  "ackTimeout": 20
}
```

3. Pinging results when tester returns 200:

```
_id: ObjectId('65ad1ac8cbcc4021111c0ff0')
jobId: "4b84825c-0b2f-499a-ac7e-059283017cb4"
timestamp: 1705843400
expiresAt: 2024-01-21T21:23:20.647+00:00
result: "SUC"
_class: "pl.mimuw.worker.entity.MonitorResultEntity"
```

```
_id: ObjectId('65ad1acdcbbcc4021111c0ff1')
jobId: "4b84825c-0b2f-499a-ac7e-059283017cb4"
timestamp: 1705843405
expiresAt: 2024-01-21T21:23:25.648+00:00
result: "SUC"
_class: "pl.mimuw.worker.entity.MonitorResultEntity"
```

4. Results after POST request to tester service (now it returns 500):

```
_id: ObjectId('65ad1bdbcbcc4021111c1027')
jobId: "4b84825c-0b2f-499a-ac7e-059283017cb4"
timestamp: 1705843675
expiresAt: 2024-01-21T21:27:55.647+00:00
result: "FAIL"
_class: "pl.mimuw.worker.entity.MonitorResultEntity"
```

```
_id: ObjectId('65ad1be0cbcc4021111c1028')
jobId: "4b84825c-0b2f-499a-ac7e-059283017cb4"
timestamp: 1705843680
expiresAt: 2024-01-21T21:28:00.647+00:00
result: "FAIL"
_class: "pl.mimuw.worker.entity.MonitorResultEntity"
```

5. We can check in the database that the scheduled buzzator added a record about last detected failure of the job, to avoid spamming alerts if the monitored service is still down. Also, it sent an event to inform the primary admin about the failure.

```
▶ _id: ObjectId('65ad1c1d5096a13117fb512b')
  jobId: "4b84825c-0b2f-499a-ac7e-059283017cb4"
  lastProcessedPingTimestamp: 1705843920
```

6. The notifier received the event about sending an alert to the admin and saved the information about the notification in the database. After that it sent an email to

*****@gmail.com:

```
_id: ObjectId('65ad1c1f2d87a256992eale3')
notificationId: "7aee0385-d147-4176-93a5-c0e4be5e12dc"
jobId: "4b84825c-0b2f-499a-ac7e-059283017cb4"
▼ primaryAdmin: Object
  email: "*****@gmail.com"
  phoneNumber: null
▼ secondaryAdmin: Object
  email: "*****@gmail.com"
  phoneNumber: null
ackTimeoutSecs: 20
acked: false
```

> i	2024-01-21 14:29:01.278 CET	POST 200 713 B 3,6 s APIs-Google; (+https://developers.goo... https://wewu-notifier-dhmti3ww3a-
> i	2024-01-21 14:29:03.114 CET	Default STARTUP TCP probe succeeded after 1 attempt for container "wewu--notifier-1" on port 8080.
▼ i	2024-01-21 14:29:03.312 CET	Escalation number
▼ {		{
		insertId: "65ad1c1f00076c5325bcc507"
▼ jsonPayload: {		{
		type: "<class 'int'>"
		escalation_number: 0
		function_name: "wewu_notifier"
		message: "Escalation number"
		}
▶ labels: {2}		{
		logName: "projects/wewu-410223/logs/run.googleapis.com%2Fstdout"
		receiveTimestamp: "2024-01-21T13:29:03.491188623Z"
▶ resource: {2}		{
		severity: "INFO"
▶ sourceLocation: {3}		{
		timestamp: "2024-01-21T13:29:03.312113Z"
		}
		}
> i	2024-01-21 14:29:04.436 CET	Publishing task to notifier topic
> i	2024-01-21 14:29:04.525 CET	Sending email to admin

[WEWU] Service fail detected Odebrane x



Wewu Alerter wewu.alert.inator@gmail.com [przez](#) bnc3.mailjet.com
do mnie ▼

[Przetłumacz na polski](#)



Fail for job 4b84825c-0b2f-499a-ac7e-059283017cb4 detected.

ACK: https://us-east1-wewu-410223.cloudfunctions.net/wewu_acker?notificationId=7aee0385-d147-4176-93a5-c0e4be5e12dc

7. After roughly 20 seconds (time for ACK specified in request) another notification was sent (attached mail from *****@gmail.com):

> i	2024-01-21 14:29:25.514 CET	POST 200 654 B 868 ms APIs-Google; (+https://developers.goo... https://wewu-notifier-dhmti3ww3a-
✓ i	2024-01-21 14:29:25.528 CET	Escalation number
<pre>{ insertId: "65ad1c350080f52a9cc4367" jsonPayload: { escalation_number: 1 message: "Escalation number" function_name: "wewu_notifier" type: "<class 'int'>" } labels: {2} logName: "projects/wewu-410223/logs/run.googleapis.com%2Fstdout" receiveTimestamp: "2024-01-21T13:29:25.531799287Z" resource: {2} severity: "INFO" sourceLocation: {3} timestamp: "2024-01-21T13:29:25.528066Z" }</pre>		
> i	2024-01-21 14:29:25.622 CET	Sending email to admin

[WEWU] Service fail detected 🔗 Odebrane x



Wewu Alert wewu.alert.inator@gmail.com [przez](#) bnc3.mailjet.com
do mnie ▾

Fail for job 4b84825c-0b2f-499a-ac7e-059283017cb4 detected.

ACK: https://us-east1-wewu-410223.cloudfunctions.net/wewu_acker?notificationId=7aee0385-d147-4176-93a5-c0e4be5e12dc

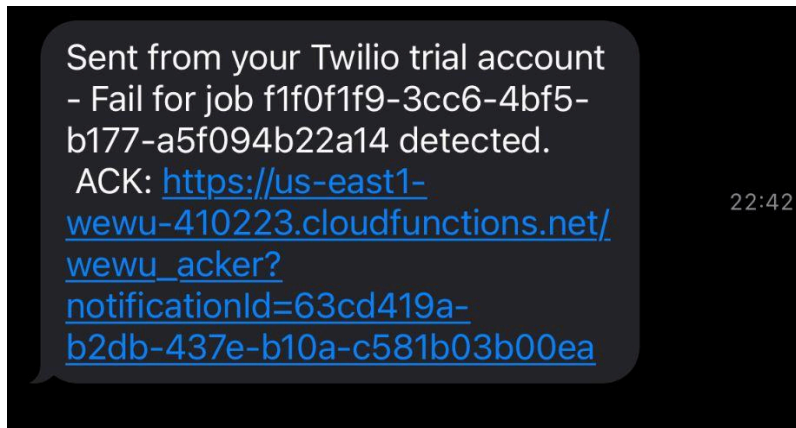
8. After another POST request to test service (now it returns 200) no more notifications arrive and service does not detect more failures (as the service is working).

Additional Notes:

- Notification can be ACKed by clicking the link in the received alert. If the first admin will ACK the message before *ackTimeoutSecs* has passed since sending the alert, the second admin won't be notified. We can observe this behavior in the logs (*ackTimeoutSecs* for this task was set to 60):

> i	2024-01-21 22:41:50.069 CET	POST 200 655 B 4,1 s APIs-Google;
> i	2024-01-21 22:41:50.089 CET	Escalation number
> i	2024-01-21 22:41:52.672 CET	Publishing task to notifier topic
> i	2024-01-21 22:42:52.749 CET	POST 200 653 B 100 ms APIs-Google;
> i	2024-01-21 22:42:52.768 CET	Escalation number
> i	2024-01-21 22:42:52.862 CET	Notification already acked, skipping

- Notification can be also send through SMS:



Load test:

Testing settings:

- Worker service can process up to 100 tasks.
- Maximum number of pods equal to 20.
- Tasks with pretty low poll frequencies (from 8s to 15s).
- WeWu-scheduler and WeWu-buzzator schedule to run every 1 minute.

Test process:


- Create an initial load of tasks equal to 200 (send requests to the API) and wait 2 min.
- Add 15 new tasks every 6 seconds (send requests to the API).

We used a python script - [link](#).

Note: We requested to monitor only services with high-availability (like google.com or facebook.com). At the end we ran the script for a while again to add some more requests, which explains the spike we can observe in the metrics. Below you can find results gathered from the metrics we find the most informative about the state of our service:


WeWu-API:

This metric reflects characteristics of our test. Spike at the beginning and at the end. Some idle time between spikes and requests coming at a consistent rate.



WeWu-scheduler:

Scheduler execution time is related to the amount of requests to the api as well as all currently monitored services. Here we can see it's execution time is longer at the beginning, later it has nothing to schedule for a while (as all tasks are processed by workers). After that its execution time lengthens after requests start coming at a consistent rate. Important difference between scheduler and api is that after the test, when no more requests are coming, it still has old jobs to schedule again.



WeWu-copy-and-paste-inator:



This is probably not the most interesting chart, but it clearly shows the scaling capabilities of Cloud Functions – when we created the first 200 messages the number of instances raised in order to match the current demand to process all the incoming events.




WeWu-worker:

In the worker, we wanted to observe two main things: the workers take on new tasks, and with the increase of the load – the number of pods increase accordingly to the specified autoscaler. On the first two charts we can clearly see the correlation between the number of monitoring tasks in the Pub/Sub and the number of pings sent per second, which indicates that workers correctly pull new tasks if possible. The creation of the new pods we could observe in the real-time, but it is also reflected on the cluster CPU/memory limits since it was scaling up in the meantime to meet the resources requirements requested by the pods.

Additionally, we can look up the GKE logs in order to check if the new pods were requested in the sensible time intervals.





Histogram




Query results ~22 log entries

SEVERITY	TIME	CET	↑	↓	SUMMARY	Edit	Summary fields	Wrap lines
Showing logs for time specified in query. To view more results update your query.								
>	i	2024-01-21 15:48:47.000			worker-service-786654bc55-jwlh6	Created	container	worker-service
>	i	2024-01-21 15:51:33.000			worker-service-786654bc55-zvhw1	Created	container	worker-service
>	i	2024-01-21 16:04:06.000			worker-service-786654bc55-n6drz	Created	container	worker-service
>	i	2024-01-21 16:05:10.000			worker-service-786654bc55-vqdk5	Created	container	worker-service
>	i	2024-01-21 16:05:10.000			worker-service-786654bc55-lmbzs	Created	container	worker-service
>	i	2024-01-21 16:05:10.000			worker-service-786654bc55-9pmb5	Created	container	worker-service
>	i	2024-01-21 16:08:19.000			worker-service-786654bc55-chdgj	Created	container	worker-service
>	i	2024-01-21 16:10:19.000			worker-service-786654bc55-blmq2	Created	container	worker-service
>	i	2024-01-21 16:13:07.000			worker-service-786654bc55-zh7hg	Created	container	worker-service
>	i	2024-01-21 16:13:07.000			worker-service-786654bc55-v58xm	Created	container	worker-service
>	i	2024-01-21 16:13:07.000			worker-service-786654bc55-t42hr	Created	container	worker-service
>	i	2024-01-21 16:15:06.000			worker-service-786654bc55-8m95f	Created	container	worker-service
>	i	2024-01-21 16:15:06.000			worker-service-786654bc55-v42cl	Created	container	worker-service
>	i	2024-01-21 16:16:23.000			worker-service-786654bc55-cm99w	Created	container	worker-service
>	i	2024-01-21 16:16:23.000			worker-service-786654bc55-skkhz	Created	container	worker-service

WeWu-buzzator:

We see that buzzator's execution time is directly related to the number of pings. After the start of the test its execution time rose consistently, even after the test ended, as the services were still monitored. We can see that even with this short period of time and only about 2k requests its execution time rose considerably and sharding would be necessary.



Stress test:

Testing settings:

- Worker service can process up to 100 tasks.
- Maximum number of pods equal to 20.
- Tasks with pretty high poll frequencies (from 1s to 2s, it changed in the meantime).
- WeWu-scheduler and WeWu-buzzator schedule to run every 1 minute.

Test process:


- Create an initial load of tasks equal to 100 (send requests to the API) and wait 2 min.
- Create a POST request to test service so it will return 500 on GET requests.
- Add 5 new tasks every 1 second, ending up with a load of 300 tasks.
- Create another load of 200. Make 5 iterations of this step running script manually in irregular intervals (up to max. a few minutes).

We used a python script - [link](#).

WeWu-worker:


This time we can see that resources used by the workers were much closer to the cluster limit than in the load test. CPU usage was higher as the services had higher frequency so

workers had to make more requests and spent less time idle.



WeWu-notifier:

Increase in the number of instances of wewu_notifier was expected as there were a lot of failures and it needed to send alerts to users. Number of instances in this diagram is a little bit exaggerated as the free mail api we are using has a limit of sending notification in a time period, so the cloud function had to rerun because of unsuccessful invocation (and also handle the new alerts at the same time).



WeWu-buzzator:

Buzzator struggled a bit with a big amount of data from pinging the services so its execution time constantly rose, just like we expected.

